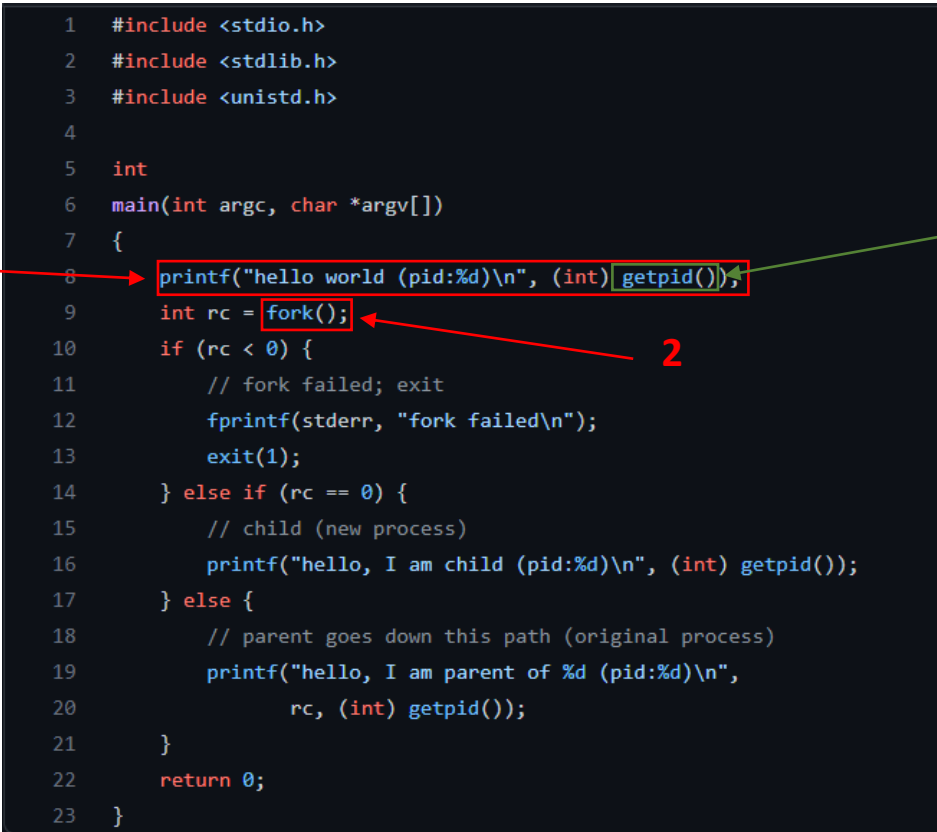


## Chapter 5: Process API

### 5.1 The `fork()` System Call


**`fork()`**: a system call is used to create a new process.



The image shows a C program snippet with line numbers 1 through 23. A red arrow labeled '1' points to line 8, which contains a `printf` statement. A green arrow labeled '1.a' points to the `getpid()` function call within the `printf` statement. A red arrow labeled '2' points to line 9, which contains the `fork()` system call. The code is as follows:

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4
5  int
6  main(int argc, char *argv[])
7  {
8      printf("hello world (pid:%d)\n", (int) getpid());
9      int rc = fork();
10     if (rc < 0) {
11         // fork failed; exit
12         fprintf(stderr, "fork failed\n");
13         exit(1);
14     } else if (rc == 0) {
15         // child (new process)
16         printf("hello, I am child (pid:%d)\n", (int) getpid());
17     } else {
18         // parent goes down this path (original process)
19         printf("hello, I am parent of %d (pid:%d)\n",
20               rc, (int) getpid());
21     }
22     return 0;
23 }
```

- 1- When it first started running, the process prints out a hello world message. (Line: 8)
  - a. Included in that message is its **process identifier**, also known as a **PID**.
    - **Process Identifier (PID)**: a unique number that is used to name the process.
      - Used to do something with the process, such as, for example, stop it from running.
- 2- The process calls the `fork()` system call, to create a new process. (Line: 9)
  - **Notes:**
    - a. The process that is created is an **(almost) exact copy** of the calling process.
    - b. Now there are two copies of the program p1 running, and both are about to return from the `fork()` System call.
      - I. **Notice:** the child isn't an exact copy. It now has its **own copy of the address space (i.e., its own private memory)**, its own registers, its own PC, and so forth.
    - c. The newly created process (called the child) doesn't start running at `main()`, rather, it just comes into life as if it had called `fork()` itself.
    - d. The value it returns to the caller of `fork()` is **different**.
      - I. The parent receives the PID of the newly created child.
      - II. The child receives a return code of zero.



You might also have noticed: **the output (of p1.c) is not deterministic.**

- ⇒ When the child process is created, there are now ***two active processes*** in the system.
  - Assuming we are running on a system with a single CPU (for simplicity), then either the child or the parent might run at that point.

```
prompt> ./p1
hello world (pid:29146)
hello, I am child (pid:29147)
hello, I am parent of 29147 (pid:29146)
prompt>
```

```
prompt> ./p1
hello world (pid:29146)
hello, I am parent of 29147 (pid:29146)
hello, I am child (pid:29147)
prompt>
```

## 5.2 The `wait()` System Call

So far, we haven't done much: just created a child that prints out a message and exits.

Sometimes, it is quite useful for a parent to wait for a child process to finish what it has been doing.

- This task is accomplished with `wait()` system call, or its more complete sibling `waitpid()`

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <sys/wait.h>
5
6  int
7  main(int argc, char *argv[])
8  {
9      printf("hello world (pid:%d)\n", (int) getpid());
10     int rc = fork();
11     if (rc < 0) {
12         // fork failed; exit
13         fprintf(stderr, "fork failed\n");
14         exit(1);
15     } else if (rc == 0) {
16         // child (new process)
17         printf("hello, I am child (pid:%d)\n", (int) getpid());
18         sleep(1);
19     } else {
20         // parent goes down this path (original process)
21         int wc = wait(NULL);
22         printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
23               rc, wc, (int) getpid());
24     }
25     return 0;
26 }
```

The parent process calls `wait()` to delay its execution until the child finishes executing.

When the child is done, `wait()` returns to the parent.

Now that you have thought a bit, here is the output:

```
prompt> ./p2
hello world (pid:29266)
hello, I am child (pid:29267) Always Child
hello, I am parent of 29267 (rc_wait:29267) (pid:29266)
prompt>
```

### 5.3 Finally, The `exec()` System Call

**`exec()`**: a system call that it **transforms** the currently running program into a different running program.


- It **does not** create a new process!

#### What it does? How `exec()` works?

```
7  int
8  main(int argc, char *argv[])
9  {
10     printf("hello world (pid:%d)\n", (int) getpid());
11     int rc = fork();
12     if (rc < 0) {
13         // fork failed; exit
14         fprintf(stderr, "fork failed\n");
15         exit(1);
16     } else if (rc == 0) {
17         // child (new process)
18         printf("hello, I am child (pid:%d)\n", (int) getpid());
19         char *myargs[3];
20         myargs[0] = strdup("wc"); // program: "wc" (word count)
21         myargs[1] = strdup("p3.c"); // argument: file to count
22         myargs[2] = NULL; // marks end of array
23         execvp(myargs[0], myargs); // runs word count
24         printf("this shouldn't print out");
25     } else {
26         // parent goes down this path (original process)
27         int wc = wait(NULL);
28         printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
29               rc, wc, (int) getpid());
30     }
31     return 0;
32 }
```


- 1- In this example, the child process calls `execvp()` in order to run the program `wc`.
  - **`wc`**: is the word counting program, which tells us how many **lines**, **words**, and **bytes** are found in the file.
- 2- In this example, the child process runs `wc` on the source file `p3.c`
  - a. It **loads code** (and static data) from that executable (**program binaries**).
  - b. **Overwrites** its current code segment (and current static data) with the loaded code.
  - c. The heap and stack and other parts of the memory space of the program are **re-initialized**.
- 3- Then the OS simply **runs** that program, **passing in any arguments as the `argv`** of that process.
  - ⇒ After the `exec()` in the child, it is almost as if `p3.c` never ran.
  - ⇒ a successful call to `exec()` never returns.

Memory before fork()




*fork()*

Memory after fork()



*exec()*

Memory before exec()



2.c

Memory after exec()



2.a & 2.b

## 5.4 Why? Motivating The API

Well, as it turns out, the separation of *fork()* and *exec()* is essential in building a UNIX shell, because it lets the shell run code after the call to *fork()* but before the call to *exec()*.

- ⇒ The separation of *fork()* and *exec()* allows the shell to do a whole bunch of useful things rather easily.

### Example 1:

The shell is just a user program. It shows you a prompt and then waits for you to type something into it. You then type a command (i.e., the name of an executable program, plus any arguments) into it.

#### Then the shell:

- 1- Figures out *where* in the file system *the executable resides (your command)*.
- 2- Calls *fork()* to create a new child process to run the command.
- 3- Calls some variant of *exec()* to run the command.
- 4- Waits for the command to complete by calling *wait()*.
- 5- The shell returns from *wait()* and prints out a prompt again, ready for your next command.

### Example 2:

*prompt >> wc p3.c > newfile.txt*

In the example above, the output of the program *wc* is redirected into the output file *newfile.txt*

- The greater-than sign is how said **redirection** is indicated.

#### The way the shell accomplishes this task is quite simple:

When the child is created (*by fork*), before calling *exec()*:

- 1- The shell closes standard output.
- 2- Opens the file *newfile.txt*.

- ⇒ Any output from the soon-to-be-running program *wc* are sent to the file instead of the screen.

Figure 5.4 (page 8) shows a program that does exactly this.

```
9  int
10 main(int argc, char *argv[])
11 {
12     int rc = fork();
13     if (rc < 0) {
14         // fork failed; exit
15         fprintf(stderr, "fork failed\n");
16         exit(1);
17     } else if (rc == 0) {
18         // child: redirect standard output to a file
19         close(STDOUT_FILENO);
20         open("./p4.output", O_CREAT|O_WRONLY|O_TRUNC, S_IRWXU);
21
22         // now exec "wc"...
23         char *myargs[3];
24         myargs[0] = strdup("wc"); // program: "wc" (word count)
25         myargs[1] = strdup("p4.c"); // argument: file to count
26         myargs[2] = NULL; // marks end of array
27         execvp(myargs[0], myargs); // runs word count
28     } else {
29         // parent goes down this path (original process)
30         int wc = wait(NULL);
31         assert(wc >= 0);
32     }
33     return 0;
34 }
```

The reason this redirection works is due to an assumption about how the operating system manages *file descriptors*.

- 1- Specifically, UNIX systems start looking for *free file descriptors at zero*.
  - In this case, **STDOUT\_FILENO** will be the first available one and thus get assigned when `open()` is called.
- 2- Subsequent writes by the child process to the standard output file descriptor, for example by routines such as **printf()**, will then be routed transparently to the newly opened file *instead of the screen*.

Here is the output of running the `p4.c` program:

```
prompt> ./p4
prompt> cat p4.output
      32      109      846 p4.c
prompt>
```

#### Notes:

- 1- When `p4` is run, *it looks as if nothing has happened*; you don't see any output printed to the screen because it has been redirected to the file **p4.output**.
- 2- You can see that when we **cat** the output file, all the expected output from running **wc** is found.

UNIX pipes are implemented in a similar way, but with the *pipe()* system call.

- The **output of one process** is connected to an in-kernel pipe (i.e., queue)
- The **input of another process** is connected to that same pipe.

- ⇒ Thus, the output of one process seamlessly is used as input to the next.
- long and useful chains of commands can be strung together.

Example:

*grep -o foo file | wc -l*

Consider looking for a word in a file, and then counting how many times said word occurs; with pipes and the utilities *grep* and *wc*, it is easy.

Topics to be covered later:

- The CPU **scheduler**, a topic we'll discuss in great detail soon, determines which process runs at a given moment in time.
- This **non-determinism**, as it turns out, leads to some interesting problems, particularly in **multi-threaded programs**; hence, we'll see a lot more non-determinism when we study **concurrency** in the second part of the book.